

Rajesh Kumar

Full Stack Java Developer & AI/GenAI Engineer | LLM Applications · RAG · Multi-Agent Systems

Chicago, IL | (773) 916-1071 | rajeshdevandla11@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

Full Stack Java Developer and AI Engineer with 5+ years of experience building production-grade backend systems and GenAI applications on AWS. Specialized in designing and shipping LLM-powered products using Amazon Bedrock, RAG pipelines, multi-agent workflows, and conversational AI — not model training, but application-layer AI engineering that delivers real business outcomes. Proficient in Java 17, Spring Boot, Angular, Python, and FastAPI. Actively targeting AI Engineer and GenAI Engineer roles where full stack depth meets applied AI.

TECHNICAL SKILLS

Languages & Frameworks:	Java 11/17, Spring Boot, Spring Data JPA, Hibernate, REST APIs, Microservices, Angular, TypeScript, JavaScript, HTML/CSS, Python, FastAPI
Cloud & DevOps:	AWS (EC2, S3, Lambda, SQS), Docker, Kafka, Jenkins, GitHub Actions, CI/CD Pipelines
AI / GenAI:	Amazon Bedrock, OpenAI API, RAG Pipelines, FAISS, LLM Integration, Prompt Engineering, Multi-Agent Workflows
Databases:	PostgreSQL, MySQL, Schema Design, Stored Procedures, Query Optimization
Testing & Tools:	JUnit, Mockito, Git, Agile/Scrum

PROFESSIONAL EXPERIENCE

Full Stack Java Developer — Mutual of Omaha (Remote, Chicago IL)

2023 – Present

- Built and delivered end-to-end features across Angular, Spring Boot, and PostgreSQL using modular service-layer architecture, reducing cross-service bugs by ~30%.
- Built 10+ REST APIs using Spring Boot, Spring Data JPA, and Hibernate supporting high-traffic insurance workflows with 99.9% uptime.
- Developed Angular reactive forms and UI components for a customer-facing portal serving 50K+ policyholders.
- Reduced synchronous API load by 45% by migrating 3 critical event workflows to AWS Lambda and SQS for async processing.
- Automated build, test, and deployment cycles using Jenkins and GitHub Actions CI/CD pipelines, cutting release time from 2 days to 4 hours.
- Implemented Apache Kafka for asynchronous event-driven communication between 4 backend services, improving system decoupling and reducing message loss by 60%.
- Integrated OpenAI API with custom prompt engineering to automate internal data classification workflows, saving ~15 hours/week of manual effort.
- Shipped a GenAI internal assistant using Amazon Bedrock with multi-turn conversation chaining, enabling faster lookup across policy documents.
- Maintained 90%+ code coverage via JUnit and Mockito; led PR reviews and enforced code quality standards across a 6-person engineering team.

Backend Java Developer — MUFG Bank / PUB Bank (Hyderabad, India)

Feb 2020 – Aug 2022

- Built and maintained 15+ REST APIs using Java 11 and Spring Boot for internal banking workflows processing 10K+ daily transactions.
- Designed SQL schemas, stored procedures, and optimized PostgreSQL/MySQL queries, reducing average query execution time by 40%.
- Integrated Apache Kafka for async event processing across 3 core banking services, cutting inter-service latency by ~35% and improving overall reliability.
- Built internal operations dashboards in vanilla JS that streamlined daily workflows for 200+ operations staff, reducing manual reporting time by ~5 hours/week.

FEATURED PROJECTS

AskDocs AI github.com/Rajeshdevandla

RAG-powered PDF chatbot using Amazon Bedrock (Claude) and FAISS for semantic search with accurate page-level citations.
Handles multi-turn Q&A; on large documents with sub-2s response time.

Stack: Python, FastAPI, Streamlit, Amazon Bedrock, FAISS, RAG

AgentFlow github.com/Rajeshdevandla

Multi-agent orchestration system (Planner, Researcher, Summarizer agents) built on Amazon Bedrock without LangChain. Custom agent base class enabling extensible, modular AI workflows.

Stack: Python, FastAPI, Amazon Bedrock, Multi-Agent Architecture

SupportGPT github.com/Rajeshdevandla

AI customer support chatbot with real-time sentiment analysis and auto-escalation logic using Amazon Bedrock. Achieves 80%+ FAQ auto-resolution rate with frustrated-customer detection.

Stack: Python, FastAPI, Amazon Bedrock, Sentiment Analysis

EDUCATION

M.S. Information Systems — DePaul University, Chicago, IL	2022 – 2024
B.S. Computer Science — GITAM University, Hyderabad, India	2017 – 2020

CERTIFICATIONS

- Building GenAI Applications with AWS — CodeSignal
- Software Engineer Certificate — HackerRank